


Product of two normal random variables

Checking the distribution functions


Generating two normal random variables

```
In[1]:= n0 = 100 000; (* number of samples *)  
In[2]:= data1 = RandomVariate[NormalDistribution[0, 1], n0];  
In[3]:= data2 = RandomVariate[NormalDistribution[0, 1], n0];  
In[4]:= Show[Histogram[data1, {-4, 4, 0.05}, "ProbabilityDensity"],  
  Plot[PDF[NormalDistribution[0, 1], x],  
    {x, -4, 4}, PlotStyle -> Thick, PlotRange -> All]]
```



Generating the product of two normal random variables

```
In[5]:= data3 = data1 * data2;  
In[6]:= Show[Histogram[data3, {-3, 3, 0.02}, "ProbabilityDensity"],  
  Plot[BesselK[0, Abs[x]] /  $\pi$ , {x, -3, 3}, PlotStyle -> Thick, PlotRange -> All]]
```



P-values for normal and the product of two normals

Numerical evaluations

```
In[7]:= precision0 = 500; (* precision for computing p-values *)
```

```
In[8]:= SetPrecision0[x_] := SetPrecision[x, precision0]
```

```
In[9]:= precision1 = 20; (* precision for outputs *)
```

```
In[10]:= N1[x_] := N[x, precision1]
```

Tail probability of normal random variable

```
In[11]:= tailnormal[x_] = Integrate[PDF[NormalDistribution[0, 1], t], {t, Abs[x], ∞}]
```

```
Out[11]=
```

$$\frac{1}{2} \operatorname{Erfc}\left[\frac{\operatorname{Abs}[x]}{\sqrt{2}}\right]$$

```
In[12]:= tailnormal[SetPrecision0[10.0]]
```

```
Out[12]=
```

```
7.619853024160526065973343251599308363504033277956960578035355462896615622059\
6481703341513851828046716081630382279898298971981706143381826035228935122022\
4590607303114569218382246575889987628945983087201467435763059338028661837282\
3290499778661334681686088451113874500815807794766957498793093705124482221553\
8148064382161426086451954772967628375408702034819605320092626303479247411111\
7309810626001221936551334469548912147659706213685301067966246882993494705325\
8629736292723371230 × 10-24
```


```
In[13]:= N1[tailnormal[SetPrecision0[10.0]]]
```

```
Out[13]=
```

```
7.6198530241605260660 × 10-24
```


```
In[14]:= Plot[tailnormal[x], {x, 0, 5}]
```

```
Out[14]=
```



```
In[15]:= Plot[Log[tailnormal[x]], {x, 0, 5}]
```

```
Out[15]=
```



Tail probability of the product of two normal random variables

```
In[16]:= tailbessel[x_] =  $\frac{1}{2}$  - Integrate[BesselK[0, t], {t, 0, Abs[x]}] /  $\pi$ 
```

```
Out[16]=
```

$$\frac{1}{2} - \frac{1}{2} \text{Abs}[x] \left(\text{BesselK}[0, \text{Abs}[x]] \text{StruveL}[-1, \text{Abs}[x]] + \text{BesselK}[1, \text{Abs}[x]] \text{StruveL}[0, \text{Abs}[x]] \right)$$

```
In[17]:= tailbessel[SetPrecision0[100.0]]
```

```
Out[17]=
```

```
1.474965847260806879759028614382444053887587255979803382535937983652503630444 \
3403296523461815455830674816321133347874152820995489362592481282019196096478 \
7809915508337250448400100825340007223075316719450155710901654358484649404614 \
8522518389102791496240357489999869628390973958445448217305520271575516822633 \
3670442653373747246793253401446572064409392670807673278647254849803556667833 \
771973525217762643039078578308433862250458611542641651149100574656174377060 \times
```

```
10-45
```


```
In[18]:= N1[tailbessel[SetPrecision0[100.0]]]
```

```
Out[18]=
```

```
1.4749658472608068798 \times 10-45
```


```
In[19]:= Plot[tailbessel[x], {x, 0, 5}, PlotRange → All]
```

Out[19]=



```
In[20]:= Plot[Log[tailbessel[x]], {x, 0, 5}, PlotRange → All]
```

Out[20]=



P-values for embeddings

one-sided p-value for the elements of embeddings with dimensions d

```
In[21]:= PvalueNormal[x_, d_] := tailnormal[x Sqrt[d]]
```


one-sided p-value for the products of two elements of embeddings with dimensions d

```
In[22]:= PvalueProduct[x_, d_] := tailbessel[x d]
```

Compare


In[23]:= `Plot[{tailnormal[x], tailbessel[x]}, {x, 0, 5}, PlotRange → All]`

Out[23]=




In[24]:= `Plot[{Log[tailnormal[x]], Log[tailbessel[x]]}, {x, 0, 5}, PlotRange → All]`

Out[24]=




In[25]:= `Plot[{PvalueNormal[x, 300], PvalueProduct[x, 300]}, {x, 0, 0.1}, PlotRange → All]`

Out[25]=



```
In[26]:= Plot[{Log[PvalueNormal[x, 300]], Log[PvalueProduct[x, 300]]},
  {x, 0, 0.1}, PlotRange -> All]
```

Out[26]=



GloVe ICA

setting

```
In[27]:= filename = "data/glove_ica"; (* base name for input csv *)
```

```
In[28]:= pvnum = 10; (* number of p-values shown in this notebook *)
```

loading embeddings

```
In[29]:= embeddings = Import[FileNameJoin[{NotebookDirectory[], filename <> ".csv"}]];
```

```
In[30]:= Dimensions[embeddings]
```

Out[30]=

```
{2, 301}
```

```
In[31]:= d0 = Dimensions[embeddings][[2]] - 1 (* dimensions *)
```

Out[31]=

```
300
```

```
In[32]:= embeddings[[1 ;; 2, 1]]
```

Out[32]=

```
{ultraviolet, light}
```

```
In[33]:= embedding1 = Table[embeddings[[1]][[i + 1]], {i, 1, d0}];
```

```
In[34]:= embedding2 = Table[embeddings[[2]][[i + 1]], {i, 1, d0}];
```

```
In[35]:= product = embedding1 embedding2;
```

largest elements

```
In[36]:= Reverse[Sort[Abs[embedding1]]][[1 ;; pvnum]]
Out[36]=
{0.534778, 0.26525, 0.219918, 0.142215,
 0.140345, 0.135273, 0.123291, 0.118807, 0.112031, 0.10384}

In[37]:= Reverse[Sort[Abs[embedding2]]][[1 ;; pvnum]]
Out[37]=
{0.553637, 0.3788, 0.23018, 0.114424, 0.113606,
 0.110458, 0.108315, 0.105231, 0.103104, 0.102278}

In[38]:= Reverse[Sort[Abs[product]]][[1 ;; pvnum]]
Out[38]=
{0.296073, 0.030351, 0.0224457, 0.0140396, 0.00698963,
 0.00558748, 0.0053346, 0.00516984, 0.00512847, 0.00509834}
```

computing p-values

```
In[39]:= pvalue1 = Table[PvalueNormal[SetPrecision0[embedding1[[i]]], d0], {i, 1, d0}];
In[40]:= pvalue2 = Table[PvalueNormal[SetPrecision0[embedding2[[i]]], d0], {i, 1, d0}];
In[41]:= pvalue3 = Table[PvalueProduct[SetPrecision0[product[[i]]], d0], {i, 1, d0}];
```

truncate digits

```
In[42]:= pvalue1 = N1[pvalue1];
In[43]:= pvalue2 = N1[pvalue2];
In[44]:= pvalue3 = N1[pvalue3];
```

smallest p - values

```
In[45]:= Sort[pvalue1][[1 ;; pvnum]]
Out[45]=
{9.9727102109255499839 × 10-21, 2.1713253147956397236 × 10-6,
 0.000069741544648986716679, 0.0068846028396609071991,
 0.0075316452423573268084, 0.0095645911121902092716, 0.016361578793168615862,
 0.019804368158821912148, 0.026163923335786963767, 0.036043532415242638207}


In[46]:= Sort[pvalue2][[1 ;; pvnum]]
Out[46]=
{4.4352624838040532268 × 10-22, 2.6721118617412933276 × 10-11,
 0.000033480324678891397434, 0.023746355809461928765,
 0.024550463178974927537, 0.027862357739569101300, 0.030323093636972657604,
 0.034177578312293119273, 0.037064199102258047676, 0.038238320398020305428}
```

```
In[47]:= Sort[pvalue3][[1 ;; pvnum]]
```


```
Out[47]= {1.1188250686729021116 × 10-40, 0.000013818835165210664030,  
0.00016900066437121847150, 0.0025596168912095961888,  
0.027602292082712734552, 0.045359478539450394649, 0.049683428069175252811,  
0.052734854137595203625, 0.053532084761371115616, 0.054120792586203003736}
```

histogram of p - values


```
In[48]:= Show[Histogram[pvalue1, 10, "ProbabilityDensity"],  
PlotStyle → Thick, PlotRange → All]
```



```
In[49]:= Show[Histogram[pvalue2, 10, "ProbabilityDensity"],  
PlotStyle → Thick, PlotRange → All]
```




```
In[50]:= Show[Histogram[pvalue3, 10, "ProbabilityDensity"],
  PlotStyle → Thick, PlotRange → All]
```



output p - values

```
In[51]:= pvalues = {pvalue1, pvalue2, pvalue3};
```

```
In[52]:= Export[FileNameJoin[{NotebookDirectory[], filename <> "-pvalue.csv"}],
  pvalues, "CSV"];
```

GloVe PCA

setting

```
In[53]:= filename = "data/glove_pca"; (* base name for input csv *)
```

```
In[54]:= pvnum = 10; (* number of p-values shown in this notebook *)
```

loading embeddings

```
In[55]:= embeddings = Import[FileNameJoin[{NotebookDirectory[], filename <> ".csv"}]];
```

```
In[56]:= Dimensions[embeddings]
```

Out[56]=
{2, 301}

```
In[57]:= d0 = Dimensions[embeddings][[2]] - 1 (* dimensions *)
```

Out[57]=
300

```
In[58]:= embeddings[[1 ;; 2, 1]]
```

Out[58]=
{ultraviolet, light}

```
In[59]:= embedding1 = Table[embeddings[[1]][[i + 1]], {i, 1, d0}];
```

```
In[60]:= embedding2 = Table[embeddings[[2]][[i + 1]], {i, 1, d0}];
```

```
In[61]:= product = embedding1 embedding2;
```

largest elements

```
In[62]:= Reverse[Sort[Abs[embedding1]]][[1 ;; pvnum]]
```

```
Out[62]= {0.202355, 0.197082, 0.174948, 0.172482,
          0.157429, 0.152461, 0.150836, 0.150187, 0.145999, 0.137225}
```

```
In[63]:= Reverse[Sort[Abs[embedding2]]][[1 ;; pvnum]]
```

```
Out[63]= {0.247417, 0.215653, 0.176051, 0.175901,
          0.174256, 0.158521, 0.145853, 0.139102, 0.132596, 0.132584}
```

```
In[64]:= Reverse[Sort[Abs[product]]][[1 ;; pvnum]]
```

```
Out[64]= {0.0238298, 0.0216126, 0.0199125, 0.0193315, 0.0170272,
          0.016883, 0.0165086, 0.0164595, 0.0160765, 0.0149764}
```

computing p-values

```
In[65]:= pvalue1 = Table[PvalueNormal[SetPrecision0[embedding1[[i]]], d0], {i, 1, d0}];
```

```
In[66]:= pvalue2 = Table[PvalueNormal[SetPrecision0[embedding2[[i]]], d0], {i, 1, d0}];
```

```
In[67]:= pvalue3 = Table[PvalueProduct[SetPrecision0[product[[i]]], d0], {i, 1, d0}];
```

truncate digits

```
In[68]:= pvalue1 = N1[pvalue1];
```

```
In[69]:= pvalue2 = N1[pvalue2];
```

```
In[70]:= pvalue3 = N1[pvalue3];
```

smallest p - values

```
In[71]:= Sort[pvalue1][[1 ;; pvnum]]
```

```
Out[71]= {0.00022839928633589803818, 0.00032060091970176385135,
          0.0012220298298100365310, 0.0014064722338352894975,
          0.0031980699317959966585, 0.0041367850647829578809, 0.0044933789779931705359,
          0.0046432738151138379601, 0.0057231311587856054694, 0.0087314904314351102950}
```

```
In[72]:= Sort[pvalue2][[1 ;; pvnum]]
```


```
Out[72]= {9.1211573884210215832 × 10-6, 0.000093775995420068982650,
          0.0011469165033407148393, 0.0011568427843692526073,
          0.0012714064817999535993, 0.0030193601548346359915, 0.0057645362598260298737,
          0.0079912418588545610080, 0.010820121694564395228, 0.010825748986848787497}
```

```
In[73]:= Sort[pvalue3][[1 ;; pvnun]]
```


```
Out[73]= {0.00010872410728436849796, 0.00022054231403184558623,
          0.00038034992569077376276, 0.00045848895053047518300,
          0.00096525251477809166805, 0.0010114822384612679314, 0.0011422849040823331301,
          0.0011606722152273658583, 0.0013146980990276005761, 0.0018825981529722625625}
```

histogram of p - values

```
In[74]:= Show[Histogram[pvalue1, 10, "ProbabilityDensity"],
              PlotStyle → Thick, PlotRange → All]
```




```
In[75]:= Show[Histogram[pvalue2, 10, "ProbabilityDensity"],
              PlotStyle → Thick, PlotRange → All]
```



```
In[76]:= Show[Histogram[pvalue3, 10, "ProbabilityDensity"],
  PlotStyle → Thick, PlotRange → All]
```

Out[76]=



output p - values

```
In[77]:= pvalues = {pvalue1, pvalue2, pvalue3};
```

```
In[78]:= Export[FileNameJoin[{NotebookDirectory[], filename <> "-pvalue.csv"}],
  pvalues, "CSV"];
```

Pythia ICA

setting

```
In[79]:= filename = "data/EleutherAI-pythia-160m_ica"; (* base name for input csv *)
```

```
In[80]:= pvnum = 10; (* number of p-values shown in this notebook *)
```

loading embeddings

```
In[81]:= embeddings = Import[FileNameJoin[{NotebookDirectory[], filename <> ".csv"}]];
```

```
In[82]:= Dimensions[embeddings]
```

Out[82]=

```
{2, 769}
```

```
In[83]:= d0 = Dimensions[embeddings][[2]] - 1 (* dimensions *)
```

Out[83]=

```
768
```

```
In[84]:= embeddings[[1 ;; 2, 1]]
```

Out[84]=

```
{ultraviolet, light}
```

```
In[85]:= embedding1 = Table[embeddings[[1]][[i + 1]], {i, 1, d0}];
```

```
In[86]:= embedding2 = Table[embeddings[[2]][[i + 1]], {i, 1, d0}];
```

```
In[87]:= product = embedding1 embedding2;
```

largest elements

```
In[88]:= Reverse[Sort[Abs[embedding1]]][[1 ;; pvnum]]
```

```
Out[88]= {0.601093, 0.149198, 0.115575, 0.107548, 0.0933143,
          0.0887963, 0.0874292, 0.0872673, 0.0863517, 0.0861818}
```

```
In[89]:= Reverse[Sort[Abs[embedding2]]][[1 ;; pvnum]]
```

```
Out[89]= {0.645316, 0.150569, 0.122209, 0.116203, 0.11412,
          0.0897541, 0.087483, 0.0742343, 0.0738345, 0.0729211}
```

```
In[90]:= Reverse[Sort[Abs[product]]][[1 ;; pvnum]]
```

```
Out[90]= {0.387895, 0.0224646, 0.0101595, 0.00820432, 0.00636319,
          0.00564241, 0.00505327, 0.00499439, 0.00494187, 0.00470642}
```

computing p-values

```
In[91]:= pvalue1 = Table[PvalueNormal[SetPrecision0[embedding1[[i]]], d0], {i, 1, d0}];
```

```
In[92]:= pvalue2 = Table[PvalueNormal[SetPrecision0[embedding2[[i]]], d0], {i, 1, d0}];
```

```
In[93]:= pvalue3 = Table[PvalueProduct[SetPrecision0[product[[i]]], d0], {i, 1, d0}];
```

truncate digits

```
In[94]:= pvalue1 = N1[pvalue1];
```

```
In[95]:= pvalue2 = N1[pvalue2];
```

```
In[96]:= pvalue3 = N1[pvalue3];
```

smallest p - values

```
In[97]:= Sort[pvalue1][[1 ;; pvnum]]
```

```
Out[97]= {1.3239142906559974835 × 10-62, 0.000017771101896645251444,
          0.00068024897130484095416, 0.0014390529714906372219,
          0.0048548028188086393439, 0.0069314959585304923610, 0.0076984070204372019074,
          0.0077940149304908772344, 0.0083544973808554944126, 0.0084623453416663916636}
```

```
In[98]:= Sort[pvalue2][[1 ;; pvnum]]
```

```
Out[98]= {7.9232621737520507911 × 10-72, 0.000015051031511772283297,
          0.00035360586495795790310, 0.00064027447504858993164,
          0.00078185208595102264153, 0.0064351209499248656046, 0.0076668636102040126032,
          0.019831526361069609108, 0.020370185096259333318, 0.021647709013704370874}
```

```
In[99]:= Sort[pvalue3][[1 ;; pvnum]]
```


```
Out[99]= {9.6611899412734605431 × 10-132, 2.9852622069772292607 × 10-9,  
0.000054438766188290305294, 0.00026806005323228060740,  
0.0012265298513882333376, 0.0022408909464902746386, 0.0036823577461693833587,  
0.0038706362305930553459, 0.0040468768615758205433, 0.0049431317624192937295}
```

histogram of p - values

```
In[100]:=
```

```
Show[Histogram[pvalue1, 10, "ProbabilityDensity"],  
PlotStyle → Thick, PlotRange → All]
```


```
Out[100]=
```




```
In[101]:=
```

```
Show[Histogram[pvalue2, 10, "ProbabilityDensity"],  
PlotStyle → Thick, PlotRange → All]
```

```
Out[101]=
```



```
In[102]:= Show[Histogram[pvalue3, 10, "ProbabilityDensity"],
  PlotStyle → Thick, PlotRange → All]
```



output p - values

```
In[103]:= pvalues = {pvalue1, pvalue2, pvalue3};

In[104]:= Export[FileNameJoin[{NotebookDirectory[], filename <> "-pvalue.csv"}],
  pvalues, "CSV"];
```

Pythia PCA

setting

```
In[105]:= filename = "data/EleutherAI-pythia-160m_pca"; (* base name for input csv *)

In[106]:= pvnum = 10; (* number of p-values shown in this notebook *)
```

loading embeddings

```
In[107]:= embeddings = Import[FileNameJoin[{NotebookDirectory[], filename <> ".csv"}]];

In[108]:= Dimensions[embeddings]

Out[108]= {2, 769}

In[109]:= d0 = Dimensions[embeddings][[2]] - 1 (* dimensions *)

Out[109]= 768
```

```

In[110]:=
  embeddings[[1 ;; 2, 1]]
Out[110]=
  {ultraviolet, light}

In[111]:=
  embedding1 = Table[embeddings[[1]][[i + 1]], {i, 1, d0}];
In[112]:=
  embedding2 = Table[embeddings[[2]][[i + 1]], {i, 1, d0}];
In[113]:=
  product = embedding1 embedding2;

```

largest elements

```

In[114]:=
  Reverse[Sort[Abs[embedding1]]][[1 ;; pvnum]]
Out[114]=
  {0.125502, 0.121365, 0.0973242, 0.0957124, 0.0929038,
   0.0920286, 0.0914271, 0.0909184, 0.0898081, 0.0876936}

In[115]:=
  Reverse[Sort[Abs[embedding2]]][[1 ;; pvnum]]
Out[115]=
  {0.128636, 0.113207, 0.10435, 0.103102, 0.101709,
   0.101052, 0.0995926, 0.0989615, 0.0986276, 0.0978361}

In[116]:=
  Reverse[Sort[Abs[product]]][[1 ;; pvnum]]
Out[116]=
  {0.0156119, 0.0129396, 0.00768797, 0.00764188, 0.00667392,
   0.00666742, 0.00645442, 0.006219, 0.0061793, 0.00616578}

```

computing p-values

```

In[117]:=
  pvalue1 = Table[PvalueNormal[SetPrecision0[embedding1[[i]], d0], {i, 1, d0}];
In[118]:=
  pvalue2 = Table[PvalueNormal[SetPrecision0[embedding2[[i]], d0], {i, 1, d0}];
In[119]:=
  pvalue3 = Table[PvalueProduct[SetPrecision0[product[[i]], d0], {i, 1, d0}];

```

truncate digits

```

In[120]:=
  pvalue1 = N1[pvalue1];
In[121]:=
  pvalue2 = N1[pvalue2];
In[122]:=
  pvalue3 = N1[pvalue3];

```


smallest p - values

In[123]:=

```
Sort[pvalue1][[1 ;; pvnun]]
```

Out[123]=

```
{0.00025256038822062073907, 0.00038497381789775003943,
 0.0034970315579463093691, 0.0039953940157082905983,
 0.0050174489478795269111, 0.0053803973158651800971, 0.0056432276757008489033,
 0.0058743490944086947131, 0.0064081171035014159542, 0.0075445040780099606809}
```

In[124]:=

```
Sort[pvalue2][[1 ;; pvnun]]
```

Out[124]=

```
{0.00018202781334466191715, 0.00085256834212590114958,
 0.0019149984644905975316, 0.0021366729418871678656,
 0.0024113984950979447977, 0.0025516287425164403500, 0.0028901409804209738951,
 0.0030486483008225889579, 0.0031356652741908659164, 0.0033508684914258616032}
```

In[125]:=

```
Sort[pvalue3][[1 ;; pvnun]]
```

Out[125]=


```
{6.8193423690938106586 × 10-7, 5.7812168629015303849 × 10-6,
 0.00040969992413534862805, 0.00042554566707283324162,
 0.00094725580361946274880, 0.00095238527597095750162,
 0.0011368312340156486596, 0.0013831515735792064068,
 0.0014297337542016547266, 0.0014459480832412269150}
```

histogram of p - values

In[126]:=

```
Show[Histogram[pvalue1, 10, "ProbabilityDensity"],
  PlotStyle → Thick, PlotRange → All]
```


Out[126]=



In[127]:=

```
Show[Histogram[pvalue2, 10, "ProbabilityDensity"],
      PlotStyle → Thick, PlotRange → All]
```

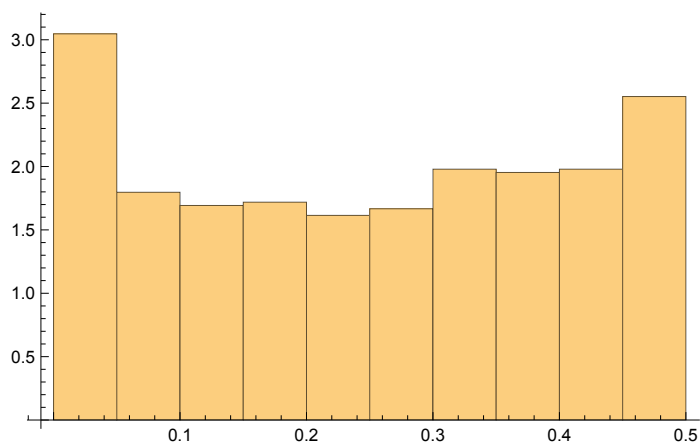
Out[127]=



In[128]:=

```
Show[Histogram[pvalue3, 10, "ProbabilityDensity"],
      PlotStyle → Thick, PlotRange → All]
```

Out[128]=



output p - values

In[129]:=

```
pvalues = {pvalue1, pvalue2, pvalue3};
```

In[130]:=

```
Export[FileNameJoin[{NotebookDirectory[], filename <> "-pvalue.csv"}],
      pvalues, "CSV"];
```

GPT2 ICA

setting

In[131]:=

```
filename = "data/gpt2_ica"; (* base name for input csv *)
```

```
In[132]:=
  pvnum = 10; (* number of p-values shown in this notebook *)
```

loading embeddings

```
In[133]:=
  embeddings = Import[FileNameJoin[{NotebookDirectory[], filename <> ".csv"}]];

In[134]:=
  Dimensions[embeddings]

Out[134]=
  {2, 769}

In[135]:=
  d0 = Dimensions[embeddings][[2]] - 1 (* dimensions *)

Out[135]=
  768

In[136]:=
  embeddings[[1 ;; 2, 1]]

Out[136]=
  {ultraviolet, light}

In[137]:=
  embedding1 = Table[embeddings[[1]][[i + 1]], {i, 1, d0}];

In[138]:=
  embedding2 = Table[embeddings[[2]][[i + 1]], {i, 1, d0}];

In[139]:=
  product = embedding1 embedding2;
```

largest elements

```
In[140]:=
  Reverse[Sort[Abs[embedding1]]][[1 ;; pvnum]]

Out[140]=
  {0.412892, 0.274985, 0.156069, 0.123778, 0.117139,
   0.11502, 0.102941, 0.101257, 0.0992244, 0.0991805}

In[141]:=
  Reverse[Sort[Abs[embedding2]]][[1 ;; pvnum]]

Out[141]=
  {0.449752, 0.321907, 0.144877, 0.128259, 0.122076,
   0.121777, 0.110015, 0.103546, 0.100394, 0.0971766}

In[142]:=
  Reverse[Sort[Abs[product]]][[1 ;; pvnum]]

Out[142]=
  {0.185699, 0.0885194, 0.0190056, 0.0179326, 0.0127208,
   0.0102743, 0.0099823, 0.00776384, 0.00774396, 0.00518704}
```

computing p-values

```
In[143]:=
  pvalue1 = Table[PvalueNormal[SetPrecision0[embedding1[[i]], d0], {i, 1, d0}];
```

```
In[144]:=
pvalue2 = Table[PvalueNormal[SetPrecision0[embedding2[[i]], d0], {i, 1, d0}]];

In[145]:=
pvalue3 = Table[PvalueProduct[SetPrecision0[product[[i]], d0], {i, 1, d0}]];
```

truncate digits

```
In[146]:=
pvalue1 = N1[pvalue1];

In[147]:=
pvalue2 = N1[pvalue2];

In[148]:=
pvalue3 = N1[pvalue3];
```

smallest p - values

```
In[149]:=
Sort[pvalue1][[1 ;; pvnum]]

Out[149]=
{1.2834480284074826699 × 10-30, 1.2625316958177614424 × 10-14,
 7.6225439022690687637 × 10-6, 0.00030151394852150871834,
 0.00058468433670341731773, 0.00071744930048294055338,
 0.0021669612570376461191, 0.0025070447163681984033,
 0.0029817028989625085584, 0.0029927759619287111553}
```

```
In[150]:=
Sort[pvalue2][[1 ;; pvnum]]

Out[150]=
{5.8756053956656604701 × 10-36, 2.3117635174655319148 × 10-19,
 0.000029730617483487919122, 0.00018941500365791267579,
 0.00035839769444402885142, 0.00036937558852008958569,
 0.0011486739457767337609, 0.0020552220152857773840,
 0.0026996042333217114831, 0.0035402288811217061808}
```

```
In[151]:=
Sort[pvalue3][[1 ;; pvnum]]

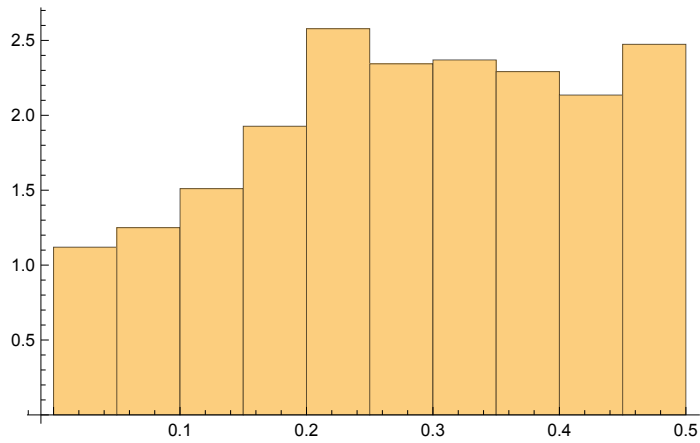
Out[151]=
{3.8397613293868146788 × 10-64, 1.4328633660188202127 × 10-31,
 4.5974783925943615532 × 10-8, 1.0766629659811503079 × 10-7,
 6.8915619827463771374 × 10-6, 0.000049599038173696981040,
 0.000062857747404859980855, 0.00038490385259725407147,
 0.00039125151305702764793, 0.0032884394819215879927}
```

histogram of p - values

In[152]:=

```
Show[Histogram[pvalue1, 10, "ProbabilityDensity"],  
      PlotStyle → Thick, PlotRange → All]
```


Out[152]=



In[153]:=

```
Show[Histogram[pvalue2, 10, "ProbabilityDensity"],  
      PlotStyle → Thick, PlotRange → All]
```


Out[153]=



In[154]:=

```
Show[Histogram[pvalue3, 10, "ProbabilityDensity"],
      PlotStyle → Thick, PlotRange → All]
```

Out[154]:=



output p - values

In[155]:=

```
pvalues = {pvalue1, pvalue2, pvalue3};
```

In[156]:=

```
Export[FileNameJoin[{NotebookDirectory[], filename <> "-pvalue.csv"}],
      pvalues, "CSV"];
```

GPT2 PCA

setting

In[157]:=

```
filename = "data/gpt2_pca"; (* base name for input csv *)
```

In[158]:=

```
pvnum = 10; (* number of p-values shown in this notebook *)
```

loading embeddings

In[159]:=

```
embeddings = Import[FileNameJoin[{NotebookDirectory[], filename <> ".csv"}]];
```

In[160]:=

```
Dimensions[embeddings]
```

Out[160]:=

```
{2, 769}
```

In[161]:=

```
d0 = Dimensions[embeddings][[2]] - 1 (* dimensions *)
```

Out[161]:=

```
768
```

```

In[162]:=
embeddings[[1 ;; 2, 1]]
Out[162]=
{ultraviolet, light}

In[163]:=
embedding1 = Table[embeddings[[1]][[i + 1]], {i, 1, d0}];

In[164]:=
embedding2 = Table[embeddings[[2]][[i + 1]], {i, 1, d0}];

In[165]:=
product = embedding1 embedding2;

```

largest elements

```

In[166]:=
Reverse[Sort[Abs[embedding1]]][[1 ;; pvnum]]
Out[166]=
{0.108743, 0.101386, 0.0995225, 0.095572, 0.0955223,
 0.0949627, 0.0922943, 0.0883422, 0.086621, 0.083538}

In[167]:=
Reverse[Sort[Abs[embedding2]]][[1 ;; pvnum]]
Out[167]=
{0.127528, 0.114267, 0.101612, 0.101204, 0.100394,
 0.0979855, 0.0977297, 0.0936943, 0.090686, 0.0897086}

In[168]:=
Reverse[Sort[Abs[product]]][[1 ;; pvnum]]
Out[168]=
{0.0095899, 0.00946041, 0.00935614, 0.00921621, 0.00715787,
 0.00666351, 0.00657745, 0.00601056, 0.00589385, 0.00566825}

```

computing p-values

```

In[169]:=
pvalue1 = Table[PvalueNormal[SetPrecision0[embedding1[[i]], d0], {i, 1, d0}];

In[170]:=
pvalue2 = Table[PvalueNormal[SetPrecision0[embedding2[[i]], d0], {i, 1, d0}];

In[171]:=
pvalue3 = Table[PvalueProduct[SetPrecision0[product[[i]], d0], {i, 1, d0}];

```

truncate digits

```

In[172]:=
pvalue1 = N1[pvalue1];

In[173]:=
pvalue2 = N1[pvalue2];

In[174]:=
pvalue3 = N1[pvalue3];

```

smallest p - values

In[175]:=

```
Sort[pvalue1][[1 ;; pvnum]]
```

Out[175]=

```
{0.0012909648877135171132, 0.0024794282531022847861,
 0.0029073704813826130227, 0.0040416670610177599312,
 0.0040581838063236048641, 0.0042481446866775299171, 0.0052678388480166798726,
 0.0071783988190064237491, 0.0081860786565549845988, 0.010304490946062993381}
```

In[176]:=

```
Sort[pvalue2][[1 ;; pvnum]]
```

Out[176]=

```
{0.00020453589951324273447, 0.00077103772939661159025,
 0.0024315389463037215391, 0.0025185733333578072587,
 0.0026994821066789520628, 0.0033092381255084655907, 0.0033807768748035388519,
 0.0047085044653103563300, 0.0059826855260951623546, 0.0064579849417334372500}
```

In[177]:=

```
Sort[pvalue3][[1 ;; pvnum]]
```

Out[177]=


```
{0.000086465316887109256062, 0.000096073535574114030709,
 0.00010458580245612786567, 0.00011721540928606733259,
 0.00063439377173455608527, 0.00095548025250574422175,
 0.0010262859085048966571, 0.0016461183196612757607,
 0.0018149479371583577933, 0.0021927956490081359983}
```

histogram of p - values

In[178]:=

```
Show[Histogram[pvalue1, 10, "ProbabilityDensity"],
  PlotStyle → Thick, PlotRange → All]
```


Out[178]=



In[179]:=

```
Show[Histogram[pvalue2, 10, "ProbabilityDensity"],
      PlotStyle → Thick, PlotRange → All]
```


Out[179]=



In[180]:=

```
Show[Histogram[pvalue3, 10, "ProbabilityDensity"],
      PlotStyle → Thick, PlotRange → All]
```

Out[180]=



output p - values

In[181]:=

```
pvalues = {pvalue1, pvalue2, pvalue3};
```

In[182]:=

```
Export[FileNameJoin[{NotebookDirectory[], filename <> "-pvalue.csv"}],
      pvalues, "CSV"];
```